

In []: Importing Necessary Libraries

In [169]: `import pandas as pd`
`import numpy as np`

In [170]: `df=pd.read_csv('CAR DETAILS FROM CAR DEKHO.csv')`

In [171]: `df`

Out[171]:

	name	year	selling_price	km_driven	fuel	seller_type	transmission	owner
0	Maruti 800 AC	2007	60000	70000	Petrol	Individual	Manual	First Owner
1	Maruti Wagon R LXI Minor	2007	135000	50000	Petrol	Individual	Manual	First Owner
2	Hyundai Verna 1.6 SX	2012	600000	100000	Diesel	Individual	Manual	First Owner
3	Datsun RediGO T Option	2017	250000	46000	Petrol	Individual	Manual	First Owner
4	Honda Amaze VX i-DTEC	2014	450000	141000	Diesel	Individual	Manual	Second Owner
...
4335	Hyundai i20 Magna 1.4 CRDi (Diesel)	2014	409999	80000	Diesel	Individual	Manual	Second Owner
4336	Hyundai i20 Magna 1.4 CRDi	2014	409999	80000	Diesel	Individual	Manual	Second Owner
4337	Maruti 800 AC BSIII	2009	110000	83000	Petrol	Individual	Manual	Second Owner
4338	Hyundai Creta 1.6 CRDi SX Option	2016	865000	90000	Diesel	Individual	Manual	First Owner
4339	Renault KWID RXT	2016	225000	40000	Petrol	Individual	Manual	First Owner

4340 rows × 8 columns

In [173]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4340 entries, 0 to 4339
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            4340 non-null   object
1   year            4340 non-null   int64
2   selling_price   4340 non-null   int64
3   km_driven       4340 non-null   int64
4   fuel            4340 non-null   object
5   seller_type     4340 non-null   object
6   transmission    4340 non-null   object
7   owner           4340 non-null   object
dtypes: int64(3), object(5)
memory usage: 271.4+ KB
```

In [174]: `df.describe()`

Out[174]:

	year	selling_price	km_driven
count	4340.000000	4.340000e+03	4340.000000
mean	2013.090783	5.041273e+05	66215.777419
std	4.215344	5.785487e+05	46644.102194
min	1992.000000	2.000000e+04	1.000000
25%	2011.000000	2.087498e+05	35000.000000
50%	2014.000000	3.500000e+05	60000.000000
75%	2016.000000	6.000000e+05	90000.000000
max	2020.000000	8.900000e+06	806599.000000

In [175]: `df['km_driven'].value_counts()`

Out[175]:

70000	236
80000	228
50000	222
120000	220
60000	215
...	
19107	1
32077	1
6480	1
118400	1
112198	1

Name: km_driven, Length: 770, dtype: int64

```
In [176]: df['year'].value_counts()
```

```
Out[176]: 2017    466
          2015    421
          2012    415
          2013    386
          2014    367
          2018    366
          2016    357
          2011    271
          2010    234
          2019    195
          2009    193
          2008    145
          2007    134
          2006    110
          2005     85
          2020     48
          2004     42
          2003     23
          2002     21
          2001     20
          1998     12
          2000     12
          1999     10
          1997      3
          1996      2
          1995      1
          1992      1
          Name: year, dtype: int64
```

```
In [178]: print(df['name'].unique())
          print(df['fuel'].unique())
          print(df['seller_type'].unique())
          print(df['transmission'].unique())
          print(df['owner'].unique())
```

```
['Maruti 800 AC' 'Maruti Wagon R LXI Minor' 'Hyundai Verna 1.6 SX' ...
'Mahindra Verito 1.5 D6 BSIII'
'Toyota Innova 2.5 VX (Diesel) 8 Seater BS IV'
'Hyundai i20 Magna 1.4 CRDi']
['Petrol' 'Diesel' 'CNG' 'LPG' 'Electric']
['Individual' 'Dealer' 'Trustmark Dealer']
['Manual' 'Automatic']
['First Owner' 'Second Owner' 'Fourth & Above Owner' 'Third Owner'
'Test Drive Car']
```

```
In [179]: data.isnull().sum()
```

```
Out[179]: name          0  
         year          0  
         km_driven     0  
         fuel          0  
         seller_type   0  
         transmission  0  
         owner         0  
         price_inlakh  0  
         dtype: int64
```

```
In [128]: # There is no null values to drop,so we can proceed further
```

```
In [185]: # Converting selling price into Lakhs  
df['price_inlakh']=df['selling_price']/100000
```

```
In [187]: df.drop(['selling_price'],axis=1,inplace=True)
```

In [191]:  df


Out[191]:

	name	year	km_driven	fuel	seller_type	transmission	owner	price_inlakh
0	Maruti 800 AC	2007	70000	Petrol	Individual	Manual	First Owner	0.60000
1	Maruti Wagon R LXI Minor	2007	50000	Petrol	Individual	Manual	First Owner	1.35000
2	Hyundai Verna 1.6 SX	2012	100000	Diesel	Individual	Manual	First Owner	6.00000
3	Datsun RediGO T Option	2017	46000	Petrol	Individual	Manual	First Owner	2.50000
4	Honda Amaze VX i-DTEC	2014	141000	Diesel	Individual	Manual	Second Owner	4.50000
...
4335	Hyundai i20 Magna 1.4 CRDi (Diesel)	2014	80000	Diesel	Individual	Manual	Second Owner	4.09999
4336	Hyundai i20 Magna 1.4 CRDi	2014	80000	Diesel	Individual	Manual	Second Owner	4.09999
4337	Maruti 800 AC BSIII	2009	83000	Petrol	Individual	Manual	Second Owner	1.10000
4338	Hyundai Creta 1.6 CRDi SX Option	2016	90000	Diesel	Individual	Manual	First Owner	8.65000
4339	Renault KWID RXT	2016	40000	Petrol	Individual	Manual	First Owner	2.25000

4340 rows × 8 columns


In []:  Exploratory Data Analysis(EDA)

```
In [196]: ▶ import seaborn as sns  
import matplotlib.pyplot as plt  
sns.displot(df['price_inlakh'])  
plt.show()
```




```
In [197]: ▶ import warnings  
warnings.filterwarnings('ignore')
```


```
In [198]: ▶ import matplotlib.pyplot as plt  
sns.boxplot(df['price_inlakh'])  
plt.show()
```




```
In [199]: ▶ sns.boxplot(data['km_driven'])  
plt.show()
```




```
In [201]: ▶ sns.countplot(df['seller_type'])  
plt.show()
```

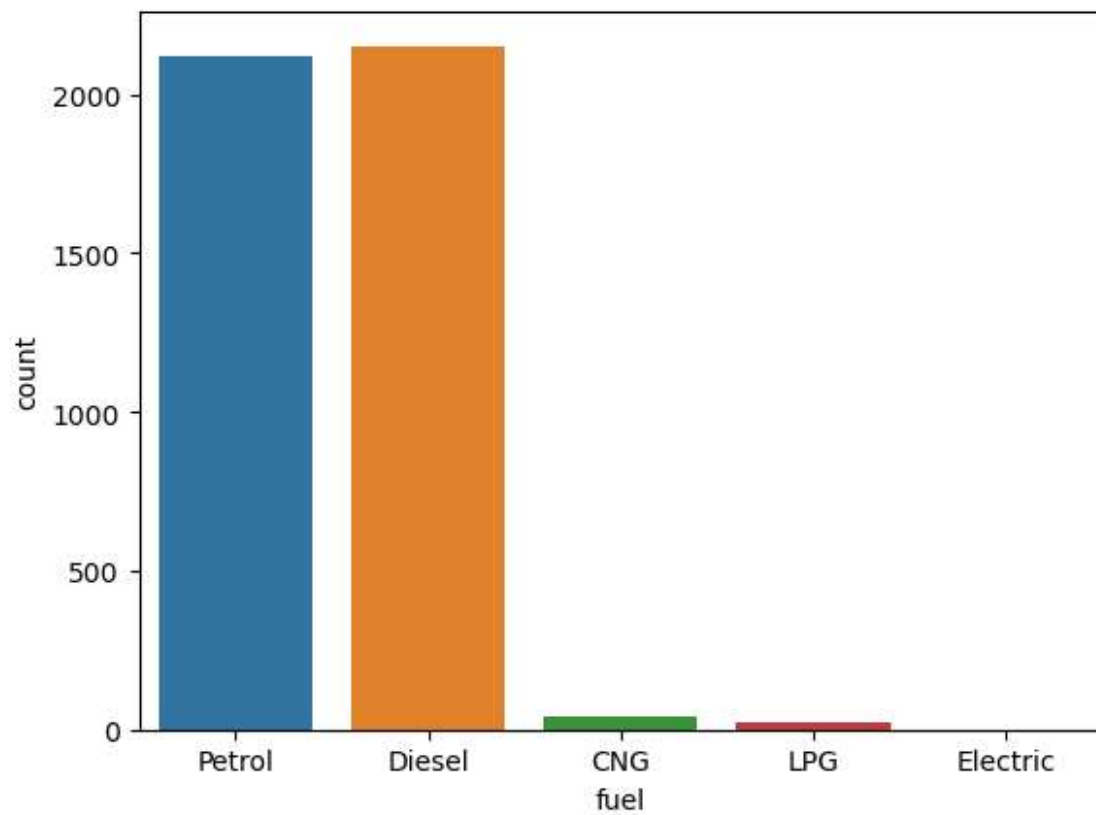


```
In [203]: ▶ sns.countplot(df['owner'])  
plt.xticks(rotation=90)  
plt.show()
```

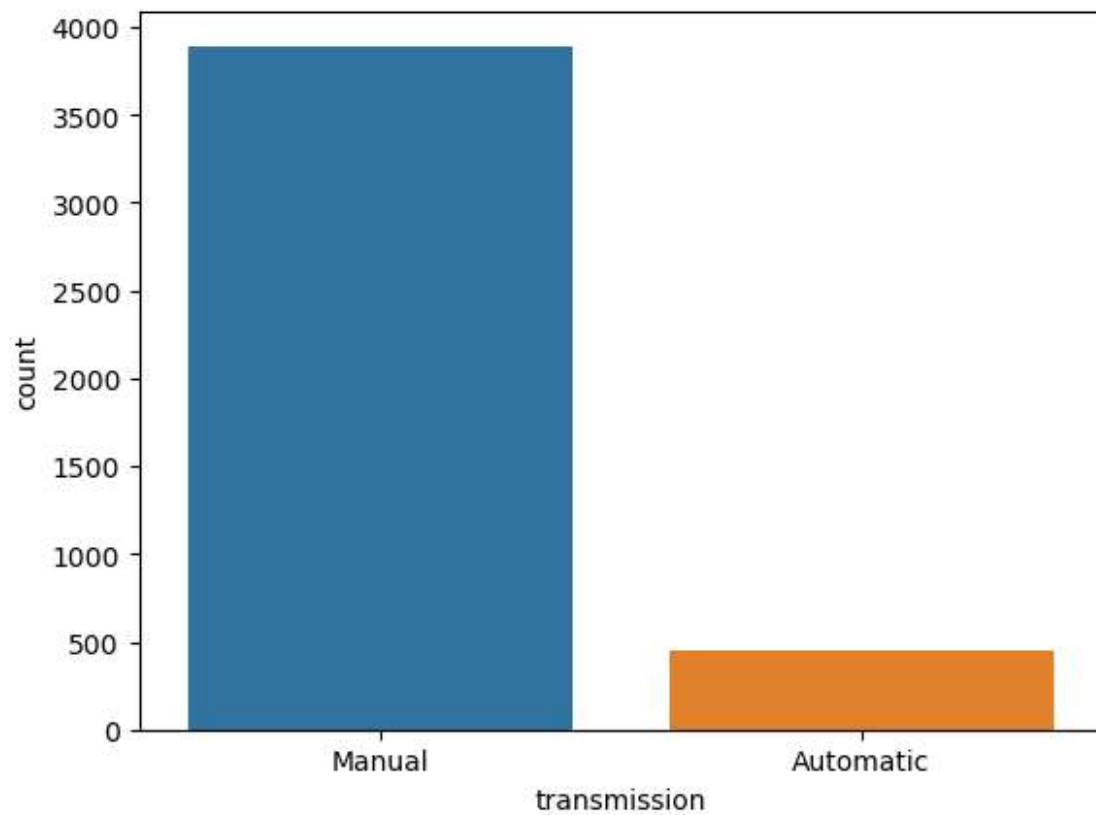


```
In [135]: ▶ # First and Second owner cars mostly sold when compared to more owners
```


```
In [205]: ▶ sns.countplot(df['fuel'])  
plt.show()
```



```
In [206]: ▶ sns.countplot(df['transmission'])  
plt.show()
```

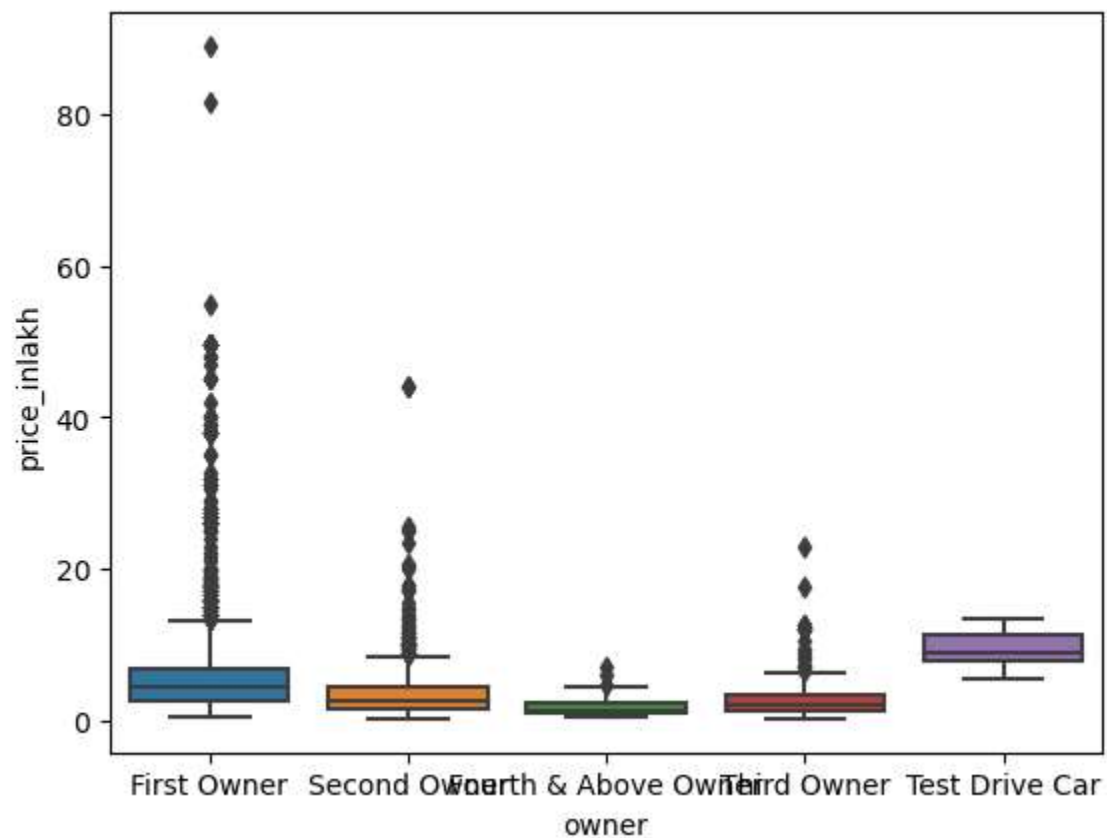


```
In [207]: ▶ corr=df.corr()  
sns.heatmap(corr)  
plt.show()
```

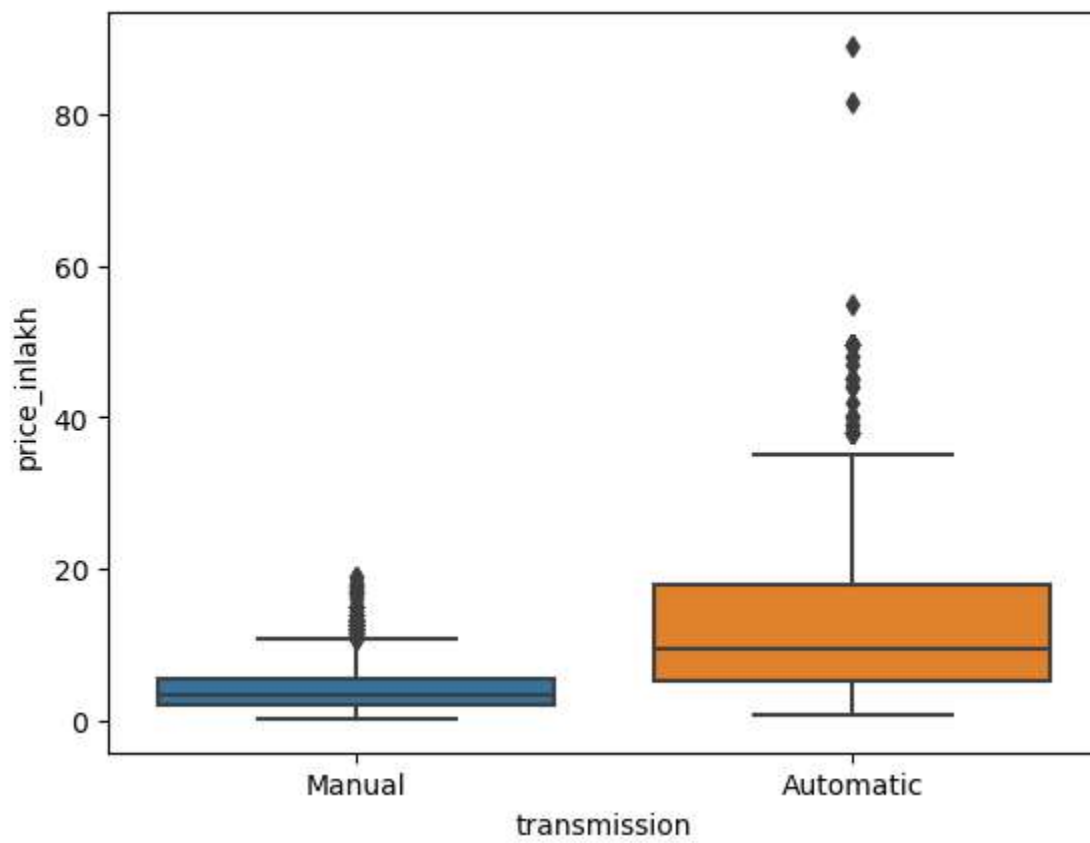


```
In [139]: ▶ # There is no such relation between the features
```

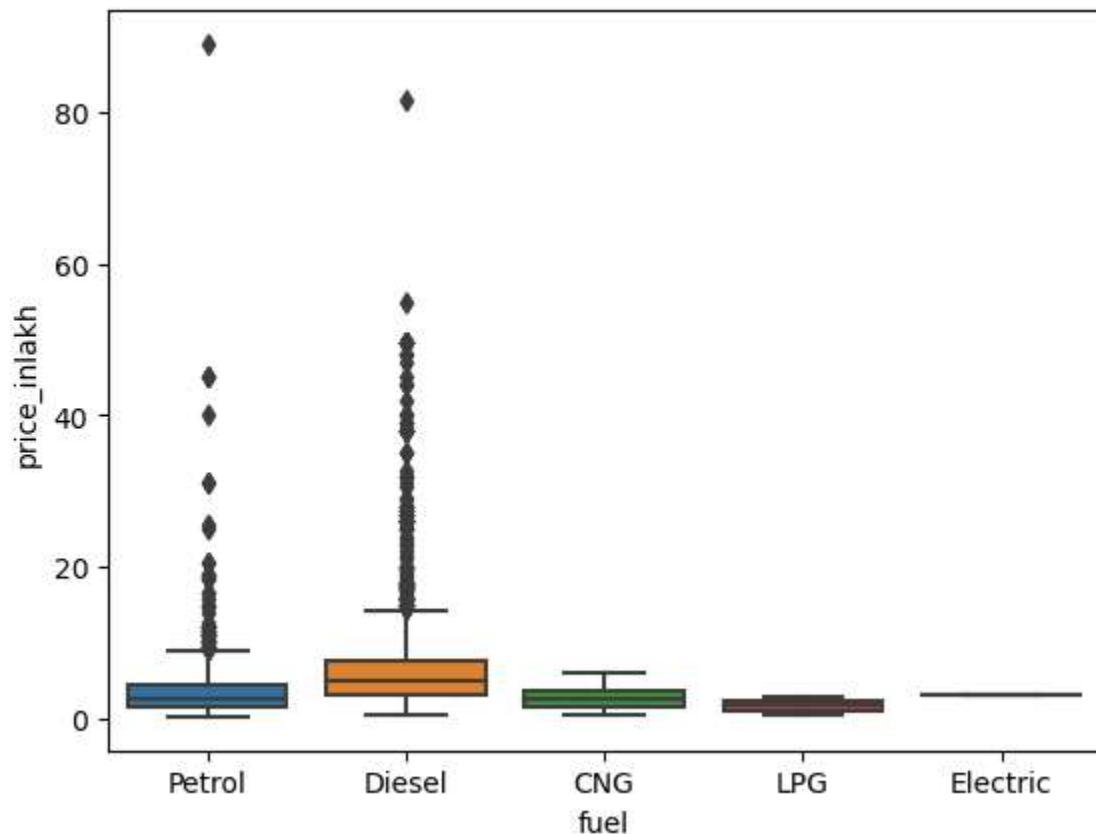
```
In [208]: sns.boxplot(y='price_inlakh',x='owner',data=df)  
plt.show()
```



```
In [210]: ▶ sns.boxplot(y='price_inlakh',x='transmission',data=df)  
plt.show()
```



```
In [146]: ▶ sns.boxplot(y='price_inlakh',x='fuel',data=data)
plt.show()
```



```
In [ ]: ▶ # After Looking transmission,fuel,owner vs price_in Lakh and boxplot of pr
# above forty Lakh.We can remove that outliers
```

```
In [211]: ▶ df2=df[df['price_inlakh']<40]
```

```
In [215]: ▶ df2=df2[df2['km_driven']<300000]
```

```
In [216]: ▶ df2.shape
```

```
Out[216]: (4301, 8)
```

```
In [ ]: ▶ # 39 records identified as outliers and removed
```

```
In [ ]: ▶ # Data Preprocessing
```

```
In [217]: ▶ from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```



```
In [218]: df2.name=le.fit_transform(df2.name)
df2.fuel=le.fit_transform(df2.fuel)
df2.transmission=le.fit_transform(df2.transmission)
df2.owner=le.fit_transform(df2.owner)
df2.seller_type=le.fit_transform(df2.seller_type)
```

```
In [219]: df2
```

Out[219]:

	name	year	km_driven	fuel	seller_type	transmission	owner	price_inlakh
0	767	2007	70000	4	1	1	0	0.60000
1	1033	2007	50000	4	1	1	0	1.35000
2	500	2012	100000	1	1	1	0	6.00000
3	113	2017	46000	4	1	1	0	2.50000
4	274	2014	141000	1	1	1	2	4.50000
...
4335	597	2014	80000	1	1	1	2	4.09999
4336	596	2014	80000	1	1	1	2	4.09999
4337	769	2009	83000	4	1	1	2	1.10000
4338	376	2016	90000	1	1	1	0	8.65000
4339	1142	2016	40000	4	1	1	0	2.25000

4301 rows × 8 columns

```
In [ ]: # Model Building
```

```
In [220]: X=df2.drop('price_inlakh',axis=1)
Y=df2['price_inlakh']
```

```
In [221]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=.20,random_st
```

```
In [222]: #Linear Regressor
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[222]:

```
LinearRegression()
LinearRegression()
```

```
In [223]: ▶ from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
y_pre_lr=lr.predict(x_test)
mae = round(mean_absolute_error(y_test, y_pre_lr), 2)
mse = round(mean_squared_error(y_test, y_pre_lr), 2)
rmse = round(np.sqrt(mse), 2)
r2 = round(r2_score(y_test, y_pre_lr), 2)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("r-squared:", r2)
```

MAE: 2.1
MSE: 10.82
RMSE: 3.29
r-squared: 0.49

```
In [224]: ▶ print(lr.score(x_train,y_train))
print(lr.score(x_test,y_test))
```

0.5000870729486298
0.4913890214038892

```
In [235]: ▶ # Random Forest Regressor
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor(n_estimators=100)
rfr.fit(x_train, y_train)
y_pre_rfr= rf_reg.predict(x_test)
print("Accuracy on Traing set: ",rfr.score(x_train,y_train))
print("Accuracy on Testing set: ",rfr.score(x_test,y_test))
```

Accuracy on Traing set: 0.9782170621584568
Accuracy on Testing set: 0.822319241792626

```
In [236]: ▶ mae = round(mean_absolute_error(y_test, y_pre_rfr), 2)
mse = round(mean_squared_error(y_test, y_pre_rfr), 2)
rmse = round(np.sqrt(mse), 2)
r2 = round(r2_score(y_test, y_pre_rfr), 2)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("r-squared:", r2)
```

MAE: 0.73
MSE: 2.36
RMSE: 1.54
r-squared: 0.89

```
In [231]: ▶ # Gradient Boosting Regressor
from sklearn.ensemble import GradientBoostingRegressor
gbr = GradientBoostingRegressor()
gbr.fit(x_train, y_train)
y_pre_gbr= gbr.predict(x_test)
print("Accuracy on Traing set: ",gbr.score(x_train,y_train))
print("Accuracy on Testing set: ",gbr.score(x_test,y_test))
```

Accuracy on Traing set: 0.8509833393341351
Accuracy on Testing set: 0.7779316920916748

```
In [232]: ▶ mae = round(mean_absolute_error(y_test, y_pre_gbr), 2)
mse = round(mean_squared_error(y_test, y_pre_gbr), 2)
rmse = round(np.sqrt(mse), 2)
r2 = round(r2_score(y_test, y_pre_gbr), 2)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("r-squared:", r2)
```

MAE: 1.25
MSE: 4.73
RMSE: 2.17
r-squared: 0.78

```
In [233]: ▶ # Decision Tree Regressor
from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor()
dtr.fit(x_train, y_train)
y_preddtr= dtr.predict(x_test)
print("Accuracy on Traing set: ",dtr.score(x_train,y_train))
print("Accuracy on Testing set: ",dtr.score(x_test,y_test))
```

Accuracy on Traing set: 0.9998629923781052
Accuracy on Testing set: 0.7341216130779781

```
In [234]: ▶ mae = round(mean_absolute_error(y_test, y_preddtr), 2)
mse = round(mean_squared_error(y_test, y_preddtr), 2)
rmse = round(np.sqrt(mse), 2)
r2 = round(r2_score(y_test, y_preddtr), 2)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("r-squared:", r2)
```

MAE: 1.12
MSE: 5.66
RMSE: 2.38
r-squared: 0.73

```
In [237]: # AdaBoost Regressor
from sklearn.ensemble import AdaBoostRegressor
abr = AdaBoostRegressor()
abr.fit(x_train, y_train)
y_pre_abr = abr.predict(x_test)
print("Accuracy on Traing set: ",abr.score(x_train,y_train))
print("Accuracy on Testing set: ",abr.score(x_test,y_test))
```

Accuracy on Traing set: 0.6200215076739344
Accuracy on Testing set: 0.5800465008927578

```
In [238]: mae = round(mean_absolute_error(y_test, y_pre_abr), 2)
mse = round(mean_squared_error(y_test, y_pre_abr), 2)
rmse = round(np.sqrt(mse), 2)
r2 = round(r2_score(y_test, y_pre_abr), 2)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("r-squared:", r2)
```

MAE: 2.14
MSE: 8.94
RMSE: 2.99
r-squared: 0.58

```
In [239]: regression_models = [lr,rfr,gbr,dtr,abr]
score_train = list()
score_test = list()

for model in regression_models :
    model.fit(x_train,y_train)
    y_pred = model.predict(x_test)

    score_train.append(model.score(x_train,y_train))
    score_test.append(model.score(x_test,y_test))
```

```
In [240]: model_names = ['Linear Regression','Random Forest Regressor','Gradient Boosting Regressor',
                        'AdaBoostRegressor']

scores = pd.DataFrame([model_names,score_train,score_test])
scores
```

Out[240]:

	0	1	2	3	4
0	Linear Regression	Random Forest Regressor	Gradient Boosting Regressor	Decision Tree Regressor	AdaBoostRegressor
1	0.500087	0.977721	0.850983	0.999863	0.614269
2	0.491389	0.813603	0.777932	0.723507	0.55971

```
In [241]: ▶ scores = scores.transpose()
scores.columns = [ 'Model','Training Set Accuracy','Testing set Accuracy']
scores
```

Out[241]:

	Model	Training Set Accuracy	Testing set Accuracy
0	Linear Regression	0.500087	0.491389
1	Random Forest Regressor	0.977721	0.813603
2	Gradient Boosting Regressor	0.850983	0.777932
3	Decision Tree Regressor	0.999863	0.723507
4	AdaBoostRegressor	0.614269	0.55971

```
In [168]: ▶ # Conclusion : Random Forest with 81 % accuracy and
           0.89 as R-Squared value has chosen for prediction
```